

```
1 1.Download the Employee Attrition Dataset
2 https://www.kaggle.com/datasets/patelprashant/employee-attrition
3 2.Perfrom Data Preprocessing
4 3.Model Building using Logistic Regression and Decision Tree and Random Forest
5 4.Calculate Performance metrics
```

```
In [4]: 1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

```
In [5]: 1 dataset = pd.read_csv("WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

```
In [6]: 1 dataset.shape
```


```
Out[6]: (1470, 35)
```

```
In [7]: 1 dataset.describe()
```

```
Out[7]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	Er
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	

8 rows × 26 columns



```
In [8]: 1 dataset.isnull().any()
```

```
Out[8]: Age                False
Attrition                 False
BusinessTravel            False
DailyRate                 False
Department                False
DistanceFromHome          False
Education                 False
EducationField            False
EmployeeCount             False
EmployeeNumber            False
EnvironmentSatisfaction   False
Gender                    False
HourlyRate                False
JobInvolvement            False
JobLevel                  False
JobRole                   False
JobSatisfaction           False
MaritalStatus             False
MonthlyIncome             False
MonthlyRate               False
```



```
In [9]: 1 #there is no null values to get the correlation we have to make sure about the null values
        2 dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1470 non-null   int64
1   Attrition                            1470 non-null   object
2   BusinessTravel                        1470 non-null   object
3   DailyRate                             1470 non-null   int64
4   Department                            1470 non-null   object
5   DistanceFromHome                      1470 non-null   int64
6   Education                             1470 non-null   int64
7   EducationField                        1470 non-null   object
8   EmployeeCount                         1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction                1470 non-null   int64
11  Gender                                 1470 non-null   object
12  HourlyRate                             1470 non-null   int64
13  JobInvolvement                        1470 non-null   int64
14  JobLevel                              1470 non-null   int64
15  JobRole                               1470 non-null   object
16  JobSatisfaction                       1470 non-null   int64
17  MaritalStatus                         1470 non-null   object
18  MonthlyIncome                         1470 non-null   int64
19  MonthlyRate                           1470 non-null   int64
20  NumCompaniesWorked                    1470 non-null   int64
21  Over18                                1470 non-null   int64
22  OverTime                              1470 non-null   int64
23  PercentSalaryHike                     1470 non-null   int64
24  PerformanceRating                     1470 non-null   int64
25  RelationshipSatisfaction                1470 non-null   int64
26  StandardHours                         1470 non-null   int64
27  StockOptionLevel                      1470 non-null   int64
28  TotalWorkingYears                     1470 non-null   int64
29  TrainingTimesLastYear                  1470 non-null   int64
30  WorkLifeBalance                        1470 non-null   int64
31  YearsAtCompany                         1470 non-null   int64
32  YearsInCurrentRole                     1470 non-null   int64
33  YearsSinceLastPromotion                 1470 non-null   int64
34  YearsWithCurrManager                    1470 non-null   int64
35  
```

```
In [10]: 1 dataset.nunique()
```

```
Out[10]: Age                                43
Attrition                                2
BusinessTravel                           3
DailyRate                               886
Department                               3
DistanceFromHome                         29
Education                                5
EducationField                           6
EmployeeCount                            1
EmployeeNumber                          1470
EnvironmentSatisfaction                   4
Gender                                    2
HourlyRate                               71
JobInvolvement                           4
JobLevel                                 5
JobRole                                  9
JobSatisfaction                           4
MaritalStatus                            3
MonthlyIncome                           1349
MonthlyRate                             1427
NumCompaniesWorked                       10
Over18                                   1
OverTime                                 2
PercentSalaryHike                        15
PerformanceRating                         2
RelationshipSatisfaction                   4
StandardHours                            1
StockOptionLevel                         4
TotalWorkingYears                        40
TrainingTimesLastYear                     7
WorkLifeBalance                          4
YearsAtCompany                           37
YearsInCurrentRole                       19
YearsSinceLastPromotion                   16
YearsWithCurrManager                      18
dtype: int64
```

```
1 Here EmployeeCount, Over18, StandardHours, EmployeeNumber we don't need these for analysis
```

```
2 As the EMPLOYEECOUNT is 1, and StandardHours, Over18 are same for every company
and we don't need EmployeeNUMBER because of indexes
```


```
In [11]: 1 dataset = dataset.drop(columns = ["EmployeeCount", "Over18", "EmployeeNumber", "Sta
```

```
In [12]: 1 dataset.head()
```

Out[12]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Science
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Science
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Othe
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Science
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medica

5 rows × 31 columns




```
1 numeric_columns =
  ['DailyRate', 'Age', 'DistanceFromHome', 'MonthlyIncome', 'MonthlyRate', 'PercentSal
  aryHike', 'TotalWorkingYears',
2     'YearsAtCompany', 'NumCompaniesWorked', 'HourlyRate',
3     'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager', 'Training
  TimesLastYear']
4
5 # Our output variable is Attrition: Which is a categorical Variable.
6 categorical_columns = ['Attrition', 'OverTime', 'BusinessTravel',
  'Department', 'Education',
7     'EducationField', 'JobSatisfaction', 'EnvironmentSatisfaction', 'WorkLifeBalance',
8     'StockOptionLevel', 'Gender', 'PerformanceRating',
  'JobInvolvement',
9     'JobLevel', 'JobRole',
  'MaritalStatus', 'RelationshipSatisfaction']
```

Data Visualization

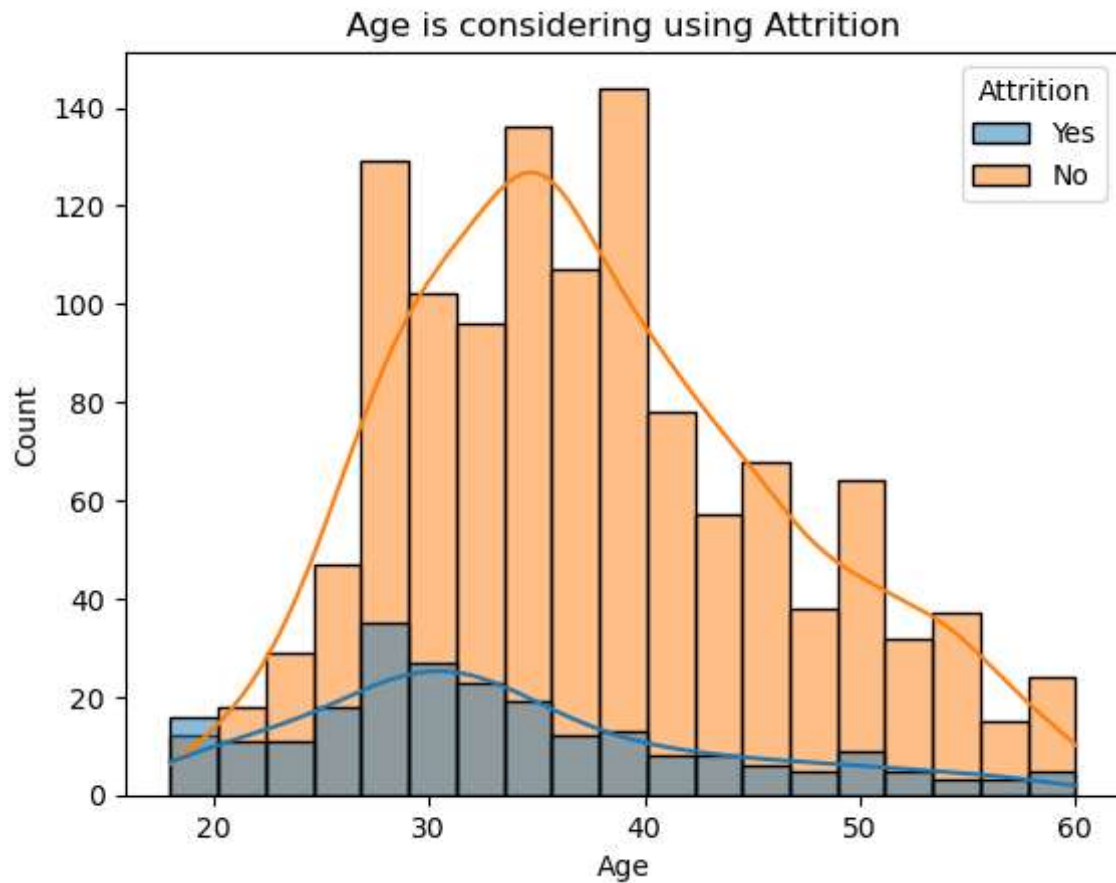
```
In [13]: 1 sns.countplot(x = "Attrition", data = dataset)
         2 plt.title("Count of Attrition")
```

```
Out[13]: Text(0.5, 1.0, 'Count of Attrition')
```



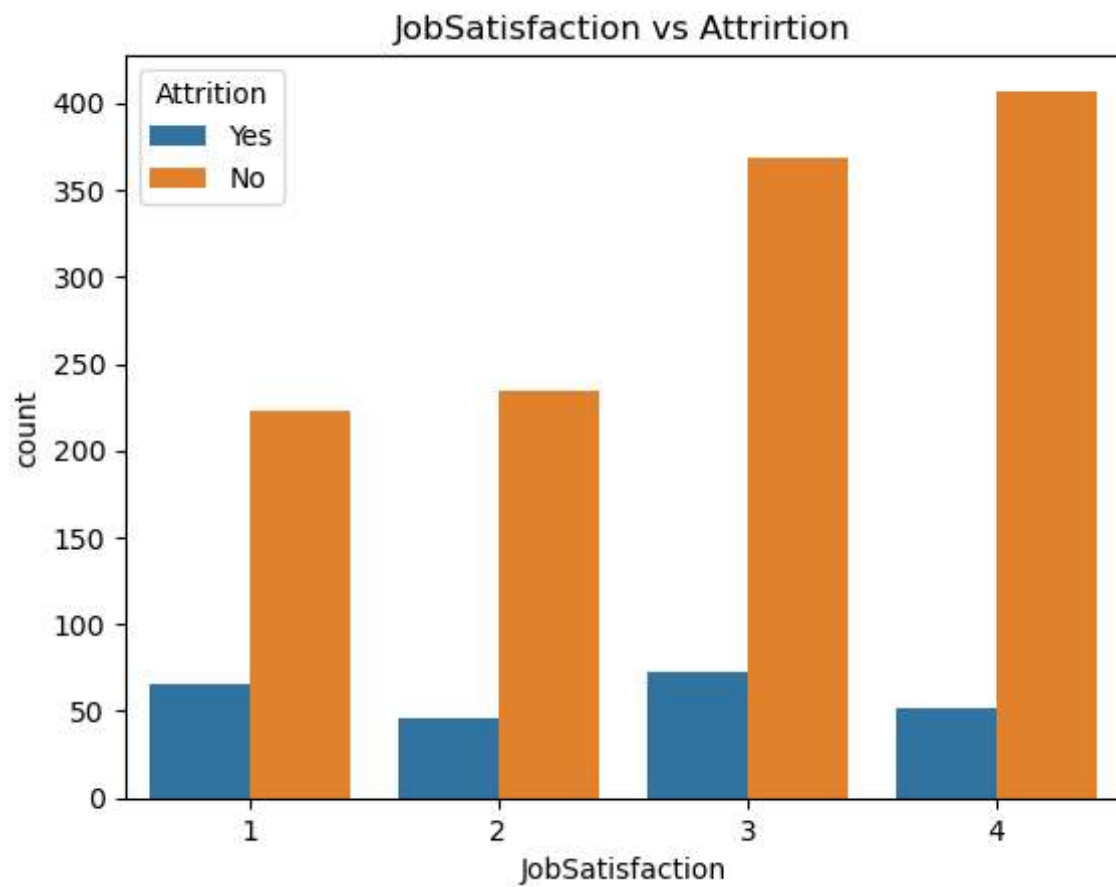
```
In [14]: 1 # checking the relation between AGE and Attrition
2 sns.histplot(data = dataset, x = 'Age', hue = 'Attrition', kde = True)
3 plt.title('Age is considering using Attrition')
```

Out[14]: Text(0.5, 1.0, 'Age is considering using Attrition')




```
In [15]: 1 # for less count variables we are using countplots instead histplots
2 sns.countplot(x = 'JobSatisfaction', hue = 'Attrition', data = dataset)
3 plt.title('JobSatisfaction vs Attrirtion')
4
```

Out[15]: Text(0.5, 1.0, 'JobSatisfaction vs Attrirtion')



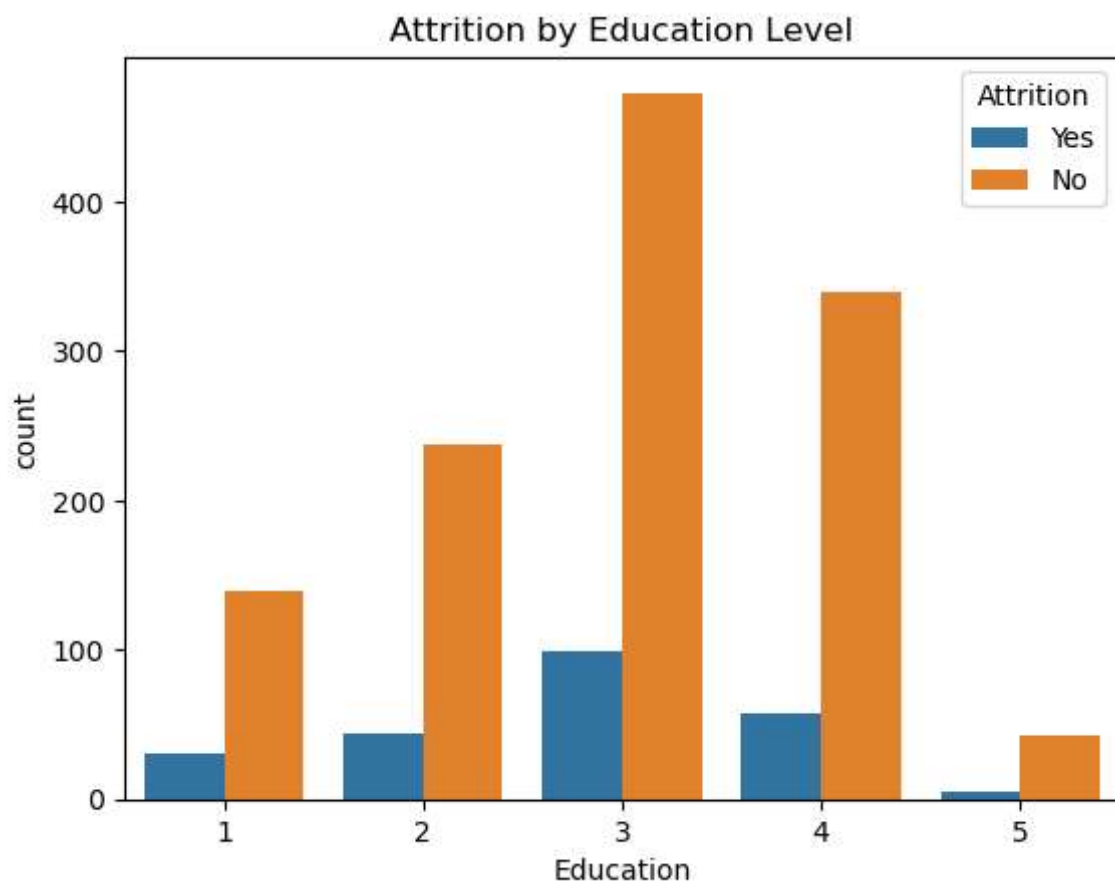
```
In [16]: 1 sns.countplot(x='Department', hue='Attrition', data=dataset)
2 plt.title('Attrition by Department')
```

Out[16]: Text(0.5, 1.0, 'Attrition by Department')




```
In [17]: 1 sns.countplot(x='Education', hue='Attrition', data=dataset)
2 plt.title('Attrition by Education Level')
```

Out[17]: Text(0.5, 1.0, 'Attrition by Education Level')




```
In [18]: 1 sns.countplot(x='WorkLifeBalance', hue='Attrition', data=dataset)
2 plt.title('Attrition by Work-Life Balance')
```

Out[18]: Text(0.5, 1.0, 'Attrition by Work-Life Balance')




```
In [19]: 1 sns.histplot(data=dataset, x='MonthlyIncome', hue='Attrition', kde=True)
2 plt.title('Distribution of Monthly Income by Attrition')
```

Out[19]: Text(0.5, 1.0, 'Distribution of Monthly Income by Attrition')



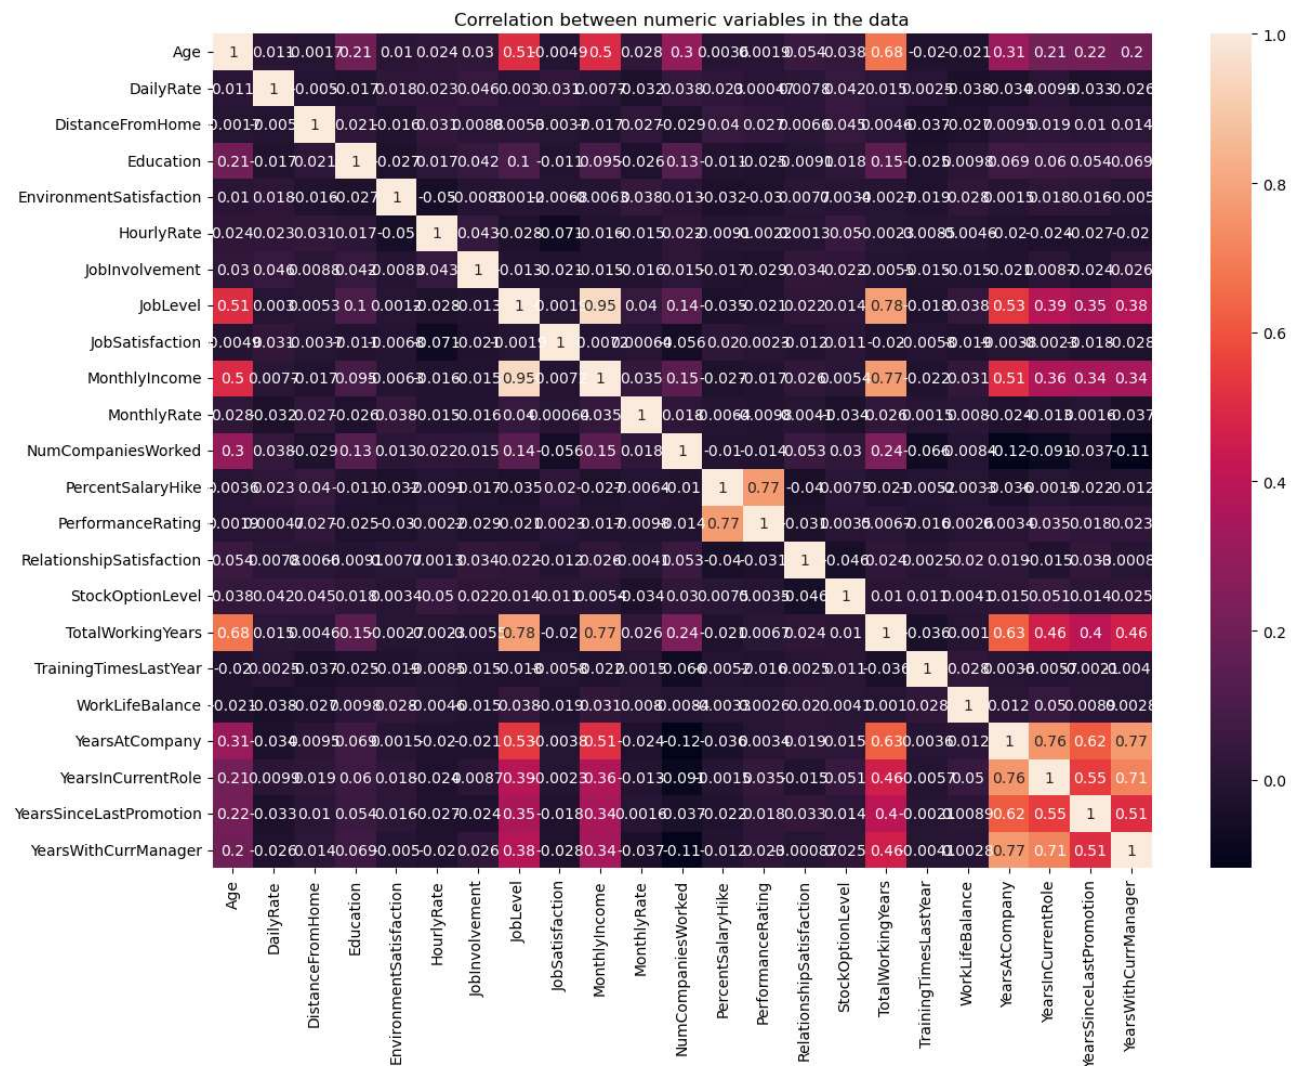

```
In [20]: 1 sns.histplot(data=dataset, x='YearsAtCompany', hue='Attrition', kde=True)
        2 plt.title('Distribution of Years at Company by Attrition')
```

Out[20]: Text(0.5, 1.0, 'Distribution of Years at Company by Attrition')



```
In [21]: 1 #Correlation Matrix (HEatmap)
2 correlation = dataset.corr(numeric_only = True)
3 plt.figure(figsize = (14,10))
4 sns.heatmap(correlation, annot =True)
5 plt.title('Correlation between numeric variables in the data')
```


Out[21]: Text(0.5, 1.0, 'Correlation between numeric variables in the data')



Checking Outliers


```
In [22]: 1 print(dataset.boxplot("Age"))  
2
```

Axes(0.125,0.11;0.775x0.77)



```
In [23]: 1 print(dataset.boxplot("DailyRate"))
```

Axes(0.125,0.11;0.775x0.77)



```
In [41]: 1 # Define the target variable (dependent variable)  
2 y = dataset['Attrition']  
3 y = pd.DataFrame(y)  
4 # Define the independent variables (features)  
5 X = dataset.drop('Attrition', axis=1)
```

```
In [42]: 1 from sklearn.preprocessing import LabelEncoder
2 le = LabelEncoder()
3 y.Attrition = le.fit_transform(y.Attrition)
```

```
In [43]: 1 # Perform one-hot encoding on categorical columns
2 X_encoded = pd.get_dummies(X, drop_first=True)
3 X_encoded.head()
```

Out[43]:

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInvolvement	J
0	41	1102	1	2	2	94	3	
1	49	279	8	1	3	61	2	
2	37	1373	2	2	4	92	2	
3	33	1392	3	4	4	56	3	
4	27	591	2	1	1	40	3	

5 rows × 44 columns

```
In [44]: 1 from sklearn.preprocessing import MinMaxScaler
2
3 # Initialize the scaler
4 scaler = MinMaxScaler()
5
6 # Fit and transform the scaled features
7 X_scaled = scaler.fit_transform(X_encoded)
8
9 # Convert the scaled features back to a DataFrame (optional)
10 X_scaled_df = pd.DataFrame(X_scaled, columns=X_encoded.columns)
```

```
In [45]: 1 X_scaled_df.head()
```

Out[45]:

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInvolvement
0	0.547619	0.715820	0.000000	0.25	0.333333	0.914286	0.666667
1	0.738095	0.126700	0.250000	0.00	0.666667	0.442857	0.333333
2	0.452381	0.909807	0.035714	0.25	1.000000	0.885714	0.333333
3	0.357143	0.923407	0.071429	0.75	1.000000	0.371429	0.666667
4	0.214286	0.350036	0.035714	0.00	0.000000	0.142857	0.666667

5 rows × 44 columns

```
In [46]: 1 from sklearn.model_selection import train_test_split
2
3 # Split the data into training and testing sets (e.g., 80% train, 20% test)
4 X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
5
```

```
In [58]: 1 from sklearn.linear_model import LogisticRegression
2 from sklearn.tree import DecisionTreeClassifier
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.metrics import accuracy_score, classification_report, roc_curve
5 import joblib # For model saving
```

```
In [59]: 1 # Initialize the models
2 logistic_model = LogisticRegression(random_state=42)
3 decision_tree_model = DecisionTreeClassifier(random_state=42)
4 random_forest_model = RandomForestClassifier(random_state=42)
```

```
In [60]: 1 # Training and testing the Logistic Regression model
2 logistic_model.fit(X_train, y_train)
3 logistic_predictions = logistic_model.predict(X_test)
4
5 # Training and testing the Decision Tree model
6 decision_tree_model.fit(X_train, y_train)
7 decision_tree_predictions = decision_tree_model.predict(X_test)
8
9 # Training and testing the Random Forest model
10 random_forest_model.fit(X_train, y_train)
11 random_forest_predictions = random_forest_model.predict(X_test)
```

C:\Users\pbalu\anaconda3\lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

C:\Users\pbalu\AppData\Local\Temp\ipykernel_3672\4153721992.py:10: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

random_forest_model.fit(X_train, y_train)

```
In [61]: 1 # Evaluation of Logistic Regression model
2 logistic_accuracy = accuracy_score(y_test, logistic_predictions)
3 logistic_report = classification_report(y_test, logistic_predictions)
4
5 print("Logistic Regression Model Accuracy:", logistic_accuracy)
6 print("Logistic Regression Model Classification Report:")
7 print(logistic_report)
```

Logistic Regression Model Accuracy: 0.891156462585034

Logistic Regression Model Classification Report:

	precision	recall	f1-score	support
0	0.91	0.97	0.94	255
1	0.67	0.36	0.47	39
accuracy			0.89	294
macro avg	0.79	0.67	0.70	294
weighted avg	0.88	0.89	0.88	294

In [62]:

```
1 # Evaluation of Decision Tree model
2 decision_tree_accuracy = accuracy_score(y_test, decision_tree_predictions)
3 decision_tree_report = classification_report(y_test, decision_tree_predictions)
4
5 print("Decision Tree Model Accuracy:", decision_tree_accuracy)
6 print("Decision Tree Model Classification Report:")
7 print(decision_tree_report)
```

Decision Tree Model Accuracy: 0.7585034013605442

Decision Tree Model Classification Report:

	precision	recall	f1-score	support
0	0.87	0.85	0.86	255
1	0.15	0.18	0.16	39
accuracy			0.76	294
macro avg	0.51	0.51	0.51	294
weighted avg	0.78	0.76	0.77	294

In [63]:

```
1 # Evaluation of Random Forest model
2 random_forest_accuracy = accuracy_score(y_test, random_forest_predictions)
3 random_forest_report = classification_report(y_test, random_forest_predictions)
4
5 print("Random Forest Model Accuracy:", random_forest_accuracy)
6 print("Random Forest Model Classification Report:")
7 print(random_forest_report)
```

Random Forest Model Accuracy: 0.8775510204081632

Random Forest Model Classification Report:

	precision	recall	f1-score	support
0	0.88	1.00	0.93	255
1	0.80	0.10	0.18	39
accuracy			0.88	294
macro avg	0.84	0.55	0.56	294
weighted avg	0.87	0.88	0.83	294

In [64]:

```
1 probability=logistic_model.predict_proba(X_test)[:,1]  
2 probability
```




```
Out[64]: array([8.20073192e-02, 4.57610775e-03, 2.84764545e-01, 2.46620134e-02,
7.47106590e-02, 3.35893969e-01, 3.76598534e-01, 4.35268844e-02,
5.75428136e-02, 1.91548056e-02, 4.82605900e-01, 5.31004390e-02,
1.09161135e-01, 7.33923982e-02, 3.59403058e-02, 1.48126811e-01,
2.65510432e-01, 1.59133872e-01, 1.84830439e-01, 4.91485926e-02,
5.54327555e-01, 5.28354516e-03, 4.93097782e-02, 3.10364306e-01,
1.34206572e-01, 1.31069154e-02, 8.22176238e-02, 1.87420735e-02,
8.79626036e-02, 4.16035868e-02, 1.45945490e-02, 3.80982342e-03,
5.70977884e-03, 3.23467257e-02, 5.58533146e-01, 9.30108837e-03,
2.02220261e-03, 2.15510371e-01, 8.00164005e-01, 2.18623190e-02,
6.17767824e-02, 9.44164078e-02, 5.97231133e-02, 5.92778741e-02,
6.64210992e-01, 1.53625484e-02, 7.50024521e-01, 4.92698546e-01,
5.65245414e-01, 5.04691723e-01, 1.65603669e-02, 1.59596535e-01,
1.54532482e-02, 1.01034186e-01, 1.11615339e-01, 1.55139363e-01,
4.06199017e-01, 4.40304306e-03, 5.83340416e-02, 1.02118488e-01,
4.17111979e-02, 4.31616734e-01, 3.50897695e-02, 2.66784979e-03,
2.12603753e-01, 4.83236708e-01, 3.01168692e-02, 1.79844614e-01,
1.39346934e-02, 1.28180635e-01, 1.92381167e-01, 1.39543844e-01,
2.77098459e-01, 9.73048198e-02, 6.29107449e-02, 1.04698197e-02,
2.27995772e-02, 9.73917066e-02, 2.45968040e-01, 1.03898276e-01,
7.71787154e-04, 1.64191596e-01, 7.93309023e-02, 1.73529996e-01,
8.19847865e-02, 6.88774428e-02, 3.50864195e-01, 2.00312867e-01,
2.39586049e-02, 5.88181809e-02, 7.85423485e-03, 7.14073544e-02,
3.61802220e-01, 3.37510187e-01, 5.38354521e-02, 8.80094441e-02,
2.77974011e-02, 1.74729416e-01, 1.80735795e-01, 1.19905041e-01,
4.15610981e-02, 2.58866778e-03, 8.21565865e-03, 1.09812302e-01,
3.01858107e-02, 1.31743313e-01, 4.95778154e-03, 3.44684822e-01,
4.23869539e-02, 2.30056581e-03, 6.61567838e-01, 6.70242509e-01,
1.53174831e-02, 4.92800309e-02, 2.16739871e-01, 1.21197060e-01,
2.70654448e-01, 1.90386686e-02, 1.64451858e-02, 2.08360629e-01,
1.71286565e-02, 3.39733400e-03, 3.00153368e-01, 3.34308649e-02,
5.15192449e-01, 4.72266207e-01, 7.33870964e-02, 4.29612893e-02,
8.26967693e-02, 1.32563434e-01, 4.63068816e-02, 8.83400334e-02,
9.40770788e-02, 1.89878961e-01, 2.26684119e-02, 2.76570923e-01,
7.48266951e-02, 1.79806960e-01, 2.79998011e-01, 2.36060318e-02,
6.77280987e-01, 8.39748658e-02, 3.29381015e-03, 7.27200357e-02,
1.08704569e-02, 7.91547766e-02, 4.83082857e-02, 7.19385088e-01,
8.97109071e-02, 6.34218963e-01, 3.09052619e-02, 3.92334894e-02,
1.09122869e-01, 6.38021680e-01, 2.44855864e-01, 2.33864819e-01,
1.14747640e-01, 1.32894006e-01, 3.88601097e-03, 3.02350626e-02,
2.45934302e-02, 4.63260000e-02, 6.45993888e-03, 4.48342428e-01,
1.64772245e-02, 2.75186604e-01, 1.87404561e-01, 2.66229919e-02,
4.36054609e-02, 1.90006272e-02, 1.45593790e-01, 2.12428561e-02,
4.74360321e-01, 4.71179087e-01, 5.76003467e-02, 3.75002572e-01,
7.20942293e-02, 7.14744191e-02, 6.82274033e-02, 6.91101869e-01,
1.28960480e-01, 8.17195383e-02, 3.68826773e-03, 5.13457865e-02,
3.47524821e-02, 6.49371473e-03, 2.99678824e-01, 5.22152996e-02,
6.38140758e-03, 1.71842384e-02, 2.56216642e-03, 1.48696943e-01,
3.05488265e-01, 3.20641554e-01, 5.93149233e-02, 7.98172093e-02,
2.85665019e-01, 2.78677873e-02, 2.10361087e-02, 2.97983450e-02,
2.97185063e-01, 3.19481683e-02, 4.45026877e-01, 9.55695626e-03,
1.92501193e-03, 1.28293658e-01, 9.81177617e-02, 2.92527467e-02,
1.53361928e-01, 2.83815840e-02, 1.68610315e-01, 2.10707430e-01,
4.06520793e-02, 2.37902858e-01, 3.71022012e-02, 3.68726932e-02,
2.18343436e-02, 6.20603212e-03, 4.46789365e-01, 2.63330802e-03,
8.18173972e-02, 4.64869204e-01, 1.26173772e-02, 8.45869142e-01,
1.32547763e-01, 3.80501709e-01, 6.91755514e-02, 3.80784774e-02,
1.04594576e-01, 6.74221165e-02, 6.17572824e-01, 1.86583352e-01,
7.60292528e-02, 1.00512888e-01, 5.42433610e-01, 1.54948350e-02,
1.36929171e-03, 2.11350434e-01, 8.61048198e-02, 1.60937255e-02,
1.30308944e-01, 6.18643161e-02, 5.95714470e-01, 4.99242439e-03,
9.79340488e-03, 3.48689548e-01, 1.82483314e-02, 6.32801339e-01,
1.89666695e-02, 3.93791265e-02, 2.81756067e-01, 6.59266369e-02,
3.20343837e-01, 1.35538343e-02, 1.47151375e-01, 4.45011797e-03,
1.35060009e-01, 1.03906895e-01, 4.50771030e-01, 3.38248450e-03,
```



```
7.59856525e-02, 7.30183735e-02, 3.94143808e-02, 1.71844121e-01,  
1.23617934e-01, 5.07344667e-02, 3.44269189e-03, 2.69923756e-02,  
3.53366535e-01, 9.79760293e-02, 8.53245153e-02, 4.94447346e-03,  
1.73729664e-01, 7.83438559e-01, 7.86236024e-02, 4.31031841e-02,  
9.99057041e-02, 2.66800338e-02, 4.38753833e-02, 6.99612307e-02,  
7.59968738e-03, 4.46879796e-01, 1.33800547e-01, 2.21784919e-02,  
2.57789402e-01, 1.65568655e-02, 2.24504799e-01, 7.10630716e-02,  
1.34460845e-02, 2.25161299e-02, 3.20533260e-02, 7.14196741e-02,  
4.44089210e-02, 9.72385427e-02])
```

```
In [65]: 1 # roc_curve  
2 fpr, tpr, threshholds = roc_curve(y_test, probability)
```

```
In [66]: 1 plt.plot(fpr, tpr)  
2 plt.xlabel('FPR')  
3 plt.ylabel('TPR')  
4 plt.title('ROC CURVE')  
5 plt.show()
```



```
In [67]: 1  
2 from sklearn.model_selection import GridSearchCV
```

```
In [ ]: 1
```


In [75]: 1 y_test

Out[75]:

Attrition	
1041	0
184	0
1222	1
67	0
220	0
...	...
567	0
560	0
945	0
522	0
651	0

In [76]:

```
1 grid_search = accuracy_score(y_test, y_pred)
2 logistic_report = classification_report(y_test, y_pred)
3
4
5 print("Logistic Regression Model Accuracy:", logistic_accuracy)
6 print("Logistic Regression Model Classification Report:")
7 print(logistic_report)
```

Logistic Regression Model Accuracy: 0.891156462585034

Logistic Regression Model Classification Report:

	precision	recall	f1-score	support
0	0.91	0.98	0.94	255
1	0.70	0.36	0.47	39
accuracy			0.89	294
macro avg	0.80	0.67	0.71	294
weighted avg	0.88	0.89	0.88	294

In [78]:

```
1 grid_search = accuracy_score(y_test, y_pred)
2 logistic_report = classification_report(y_test, y_pred)
3
4
5 print("Logistic Regression Model Accuracy:", logistic_accuracy)
6 print("Logistic Regression Model Classification Report:")
7 print(logistic_report)
```

Logistic Regression Model Accuracy: 0.891156462585034

Logistic Regression Model Classification Report:

	precision	recall	f1-score	support
0	0.91	0.98	0.94	255
1	0.70	0.36	0.47	39
accuracy			0.89	294
macro avg	0.80	0.67	0.71	294
weighted avg	0.88	0.89	0.88	294

In []:

1

In [79]:

```
1 from sklearn import metrics
2 # R- Square
3 # evaluating testing accuracy
4 print(metrics.r2_score(y_test,y_pred))
5
```

0.08355957767722466

In [83]:

```
1 import numpy as np
2 #mean squared error
3 print(metrics.mean_squared_error(y_test,y_pred))
4 # RMSE (Root Mean Square Error)
5 print(np.sqrt(metrics.mean_squared_error(y_test,y_pred)))
6 #mean absolute error
7 print(metrics.mean_absolute_error(y_test,y_pred))
```

0.1054421768707483

0.32471861183299655

0.1054421768707483

In []:

1

In []:

1